

# CLAUDE-POWERED JOB HUNT AUTOMATION

## V2 - Connector-First Architecture

Build less integration code. Use Claude as the brain, connectors as the hands,  
and Career Ops as the job-intelligence engine.

### What changed from V1

The first design used more custom Python adapters. V2 moves data acquisition and common actions to existing connectors/MCP/APIs wherever practical: Career Ops + its bundled Apify provider for jobs, Apify MCP/plugin for ad-hoc actor workflows, Google Workspace connectors for interactive email/calendar work, your existing Gmail worker for unattended background processing, and Telegram as the mobile approval layer.

Prepared: 8 September 2026

Target: Windows + PowerShell + Python + Node.js + Claude Code + MCP/connectors

# 1. The new design principle

Do not build a custom scraper, email client, browser robot, data-enrichment engine or calendar integration when a maintained connector already solves the problem. Keep custom code only where it gives you control over scoring, state, deduplication, guardrails and scheduling.

CONNECTOR-FIRST RULE

Need data/action?

|

+-- Reliable connector/MCP exists? --> USE IT

|

+-- Stable official API exists? ----> THIN ADAPTER

|

`-- Neither exists? -----> CUSTOM CODE only for this gap

Claude = reasoning/orchestration

Career Ops = job intelligence + artifacts

Connectors = acquisition/actions

Python = state + policies + schedules + retries

**Why this is better**  
You reduce maintenance, keep the codebase smaller, get structured data faster, and can swap providers later without rewriting your scoring/tailoring logic.

# 2. Final architecture

CLAUDE

Reasoning / routing

|

+-----+-----+

| | |

v v v

APIFY CAREER OPS GOOGLE WORKSPACE

web/jobs/research score/CV/cover Gmail/Calendar/Drive

| | |

+-----+-----+

|

v

THIN ORCHESTRATOR

SQLite + policies + dedupe + logs

|

+-----+-----+

| |

v v

TELEGRAM JOB TRACKER

mobile approvals outcomes/metrics

|

v

YOU

LinkedIn: profile link + prepared message only; manual send.

| Layer                 | Use                                                  | Do not rebuild                                       |
|-----------------------|------------------------------------------------------|------------------------------------------------------|
| Apify / job providers | Collect structured public web/job data               | Site-specific crawlers, pagination, browser plumbing |
| Career Ops            | Scan/evaluate jobs; CV/cover; contact/research modes | A second ATS/job evaluation framework                |

Career Ops + Apify + Claude + Gmail/Calendar + Telegram | September 2026

| Layer                     | Use                                                                    | Do not rebuild                                                |
|---------------------------|------------------------------------------------------------------------|---------------------------------------------------------------|
| Claude                    | Judgment, evidence mapping, personalization, classification            | Long-lived database/scheduler                                 |
| Google Workspace          | Interactive Gmail/Calendar/Drive actions                               | Manual copy/paste for every routine email/calendar task       |
| Existing Gmail API worker | Background inbox polling and deterministic sends after policy/approval | Replacing a working unattended worker just to use a connector |
| Telegram                  | Mobile queue and approvals                                             | A new web dashboard for MVP                                   |
| SQLite                    | Single source of truth                                                 | Multiple trackers before MVP works                            |

### 3. What Career Ops already gives you

Use Career Ops as a platform, not merely a reference repo. It already covers ATS scanning, structured evaluations, tailored CVs, cover letters, company research, contact discovery, outreach drafts, pipeline/outcome workflows and interview preparation. Its plugin system also includes a bundled Apify provider.

#### Important existing Apify support

Career Ops includes a keyed Apify provider. A portals.yml entry with provider: apify can run an Actor, map dataset fields into Career Ops, and write those jobs into the same pipeline as other sources. It is intentionally opt-in rather than auto-detected.

```
# Career Ops example: portals.yml
tracked_companies:
- name: "Indeed - MLOps Pune"
  provider: apify
  actor: <chosen-actor-id>
  input:
    position: "MLOps Engineer"
    location: "Pune, Maharashtra"
    maxItems: 25
  field_map:
    title: [positionName, title]
    url: url
    company: [company, companyName]
    location: [location, formattedLocation]

# .env
APIFY_TOKEN=your_apify_token

# Then Career Ops scanner consumes it
node scan.mjs
```

### 4. Two ways to use Apify - use both for different jobs

| Mode                            | Best for                                                    | How                                                                                  |
|---------------------------------|-------------------------------------------------------------|--------------------------------------------------------------------------------------|
| Career Ops Apify provider       | Repeatable production job feeds                             | Configure provider: apify in portals.yml and let Career Ops normalize results        |
| Apify MCP/plugin in Claude Code | Exploration, actor discovery, one-off research, prototyping | Claude searches Actors, inspects schemas, runs Actors and reads datasets through MCP |

The Apify Claude Code plugin currently bundles the Apify MCP server, an Apify routing agent and skills. Authentication is done once; Claude can then discover Actors, inspect their inputs/pricing/output, run appropriate Actors and retrieve dataset results.

1. Install/enable the Apify plugin for Claude Code using the current Apify instructions.

- 2. In Claude Code, open /mcp, enable the Apify MCP server, authenticate with Apify, and verify the connection.
- 3. Use the MCP path to discover/testing Actors; once you choose a reliable job Actor, pin its Actor ID and field mapping in Career Ops for scheduled production runs.

**Production rule**  
Do not ask Claude to rediscover a scraping Actor every scheduled run. Discovery is agentic; production should pin a reviewed Actor, its input schema, field map, cost cap and expected output fields.

## 5. Choose connectors by responsibility

| Need                                               | Preferred integration                                                       | Fallback                                                           |
|----------------------------------------------------|-----------------------------------------------------------------------------|--------------------------------------------------------------------|
| Fresh jobs from Greenhouse/Ashby/Lever/company ATS | Career Ops native scanners                                                  | Apify Actor where native scanner is insufficient                   |
| Job boards/public web extraction                   | Career Ops Apify provider                                                   | Direct Apify MCP/API                                               |
| Company/public-web research                        | Claude web research + approved Apify Actor when structured extraction helps | Career Ops deep/contacto modes                                     |
| Contact discovery                                  | Career Ops contacto + public sources                                        | Enrichment provider only when terms/data provenance are acceptable |
| Email in Claude session                            | Google Workspace Gmail connector                                            | Existing Gmail API code                                            |
| Background email watcher                           | Existing Gmail API worker                                                   | Scheduled service using Gmail API                                  |
| Calendar in Claude session                         | Google Calendar connector                                                   | Calendar API adapter                                               |
| Resume/cover files                                 | Career Ops                                                                  | Claude + deterministic PDF generator only for missing cases        |
| Mobile approvals                                   | Telegram Bot API                                                            | Simple local dashboard later                                       |

## 6. Keep interactive connectors and background workers separate

A connector is excellent when you are actively using Claude. A scheduled background job needs deterministic behavior even when no chat is open. Therefore, use the Google Workspace connector for human-in-the-loop work, but retain your existing Gmail worker for unattended polling, classification triggers and approved sends.

INTERACTIVE

You -> Claude -> Gmail/Calendar connector -> approval -> action

BACKGROUND

Task Scheduler -> Python worker -> Gmail API -> event

|

v

Claude

classify/reason

|

v

SQLite/Telegram

**This avoids a common mistake**  
Connector-first does not mean connector-only. The objective is to remove unnecessary integration code while keeping scheduled processes reliable.

## 7. Minimal workspace after simplification

```
E:\sri\career-agent\
|-- career-ops\
|-- linkedinnotification\
|-- orchestrator\
|   |-- main.py
|   |-- pipeline.py
|   |-- policies.py
|   |-- claude_bridge.py
|   `-- db.py
|-- profile\
|   |-- master_profile.yaml
|   |-- evidence_bank.yaml
|   |-- preferences.yaml
|   `-- voice_dna.md
|-- data\career_agent.db
|-- artifacts\<company>\<role>\...
|-- logs\
|-- CLAUDE.md
`-- .env
```

Notice what disappeared: custom scraper modules, custom calendar code, a custom contact crawler, a second job-normalization framework and a large UI layer. Add adapters only if a chosen connector cannot be invoked reliably in your target execution mode.

## 8. Evidence bank: the non-negotiable source of truth

Connectors can fetch more data, but they do not solve hallucinations. Every tailored resume, recruiter email and interview answer must be grounded in verified evidence about you.

```
# profile/evidence_bank.yaml
projects:
- id: ncbs_ceph
  title: Ceph storage deployment
  verified_skills: [Docker, Ceph, Linux, networking]
  verified_actions:
    - "Deployed/configured Ceph storage components"
    - "Used Docker to package/run components"
  prohibited_claims:
    - "Do not invent scale, savings or production Kubernetes use"

- id: linkedin_assistant
  title: LinkedIn AI Networking Assistant
  verified_skills: [Python, Gmail API, Telegram Bot API, SQLite]
  verified_actions:
    - "Detects accepted LinkedIn connections from Gmail"
    - "Produces human-reviewed networking/referral drafts"
```

### Rule

Claude may reorder, shorten and rephrase evidence. It must never invent a technology, metric, outcome, responsibility or experience duration.

## 9. Preferences and priority policy

```
locations:
  preferred: ["Pune", "Remote - India", "Hybrid - Pune"]
roles:
  primary: ["AI Engineer", "GenAI Engineer", "ML Engineer", "MLOps Engineer"]
  secondary: ["Data Engineer", "DevOps Engineer", "Cloud Engineer", "Platform Engineer"]
```

```

experience_max_years: 3
freshness_windows_days: [1, 7]
minimum_priority_score: 80
outreach:
  recruiter_daily_cap: 10
  employee_referral_requires_approval: true
  founder_outreach_requires_approval: true
  linkedin_send_is_manual: true
  
```

# 10. Fresh-job pipeline

```

Every 2-4 hours
|
+--> Career Ops native ATS scan
|
+--> configured Apify provider jobs
|
v
normalize -> dedupe -> verify job is live -> freshness check
|
v
Claude/Career Ops fit evaluation
|
+--> < 65 archive
+--> 65-79 store
+--> 80-84 prepare when capacity
`--> >=85 P0: act now
  
```

| Filter      | Behavior                                                                                 |
|-------------|------------------------------------------------------------------------------------------|
| Posting age | <=24h strongest; 2-7d accepted; >7d down-rank unless exceptional                         |
| Experience  | Reject obvious senior/staff/lead jobs; allow flexible JDs only when evidence supports it |
| Location    | Prefer Pune/Remote India based on your settings                                          |
| Source      | Prefer original employer/ATS URL over aggregator copy                                    |
| Duplicate   | Canonical URL + company/title + JD hash                                                  |
| Live check  | Do not build an application package for a closed/dead posting                            |

# 11. Interview Priority Score

| Signal                      | Weight | Question                                                     |
|-----------------------------|--------|--------------------------------------------------------------|
| Verified skill/evidence fit | 35     | Can your real experience support the core requirements?      |
| Freshness                   | 15     | How recently was the job posted?                             |
| Seniority                   | 15     | Is the experience band realistic?                            |
| Location/eligibility        | 10     | Can you realistically work there?                            |
| Human path                  | 15     | Can we reach recruiter/manager/employee/alumni?              |
| Posting/company quality     | 5      | Is the role live and credible?                               |
| Application friction        | 5      | Can a truthful, strong application be submitted efficiently? |

**Optimize for interviews, not keyword match**

A 90% technical match with no human path and an old posting can be lower priority than an 85% match posted today with a recruiter and an employee connection.

12. Application factory: only after the score gate

- 1. Select the closest base resume track: AI/GenAI, ML, Data, MLOps, or DevOps/Cloud.
- 2. Generate requirement-to-evidence matrix.
- 3. Tailor only claims supported by the evidence bank.
- 4. Use Career Ops to generate the ATS-friendly CV/PDF and cover letter where useful.
- 5. Save the job snapshot and every generated artifact under the company/role folder.
- 6. Do not spend Claude tokens tailoring P2/archive jobs.

```
artifacts/acme/mlops-engineer/  
|-- job.json  
|-- fit_report.json  
|-- evidence_matrix.md  
|-- Srinivas_Acme_MLOps_Resume.pdf  
|-- cover_letter.pdf  
|-- recruiter_email.txt  
|-- referral_message.txt  
`-- linkedin_message.txt
```

13. Human-path engine

For every P0/P1 job, the agent should search for the best legitimate route to a real person. Career Ops contacto is the first layer; structured public research or an approved enrichment service can be added where useful.

| Person                       | Priority    | Goal                                              |
|------------------------------|-------------|---------------------------------------------------|
| Recruiter / talent partner   | Highest     | Get application seen and confirm fit/process      |
| Hiring manager               | High        | Show concrete technical relevance to team/problem |
| Existing connection / alumni | High        | Warm advice/referral path                         |
| Future teammate              | Medium-high | Context/referral if appropriate                   |
| Founder/CTO at small company | Situational | Product-specific contribution angle               |

**LinkedIn boundary**

Use public research to identify a likely person and profile link, then generate the LinkedIn message. Keep the actual LinkedIn connection/message action manual. Career Ops itself rejects scraping of platforms that prohibit automated access.

14. Email: connector for convenience, API worker for autonomy

Claude currently supports Gmail search/read/draft/send/reply/forward through its Google Workspace connector, with approval required by default. This is ideal when you are reviewing a job interactively. Your existing Gmail API implementation remains useful for unattended workflows and thread/event tracking.

| Scenario                                   | Use                                                                |
|--------------------------------------------|--------------------------------------------------------------------|
| You are in Claude reviewing a P0 job       | Gmail connector: draft/send after approval                         |
| Scheduled worker detects a recruiter reply | Gmail API worker: record event, invoke classifier, notify Telegram |
| Approved recruiter email from Telegram     | Existing Gmail API sender: deterministic send + store              |

| Scenario                                        | Use                                                        |
|-------------------------------------------------|------------------------------------------------------------|
|                                                 | message/thread ID                                          |
| Calendar interview while in Claude              | Google Calendar connector                                  |
| Background calendar creation from parsed invite | Keep approval; then connector/API depending execution mode |

### Attachments

Keep the resume artifact in your own filesystem/Drive and validate that the correct role-specific file is attached before sending. Never infer the right file from a vague filename.

## 15. Telegram becomes the only dashboard you need for MVP

🔥 P0 JOB - 92/100  
 Acme | MLOps Engineer | Pune  
 Posted: 3h ago

Strong evidence: Docker, Python, infrastructure  
 Gap: Terraform (nice-to-have)

Human path:  
 Recruiter found ✓  
 Employee connection ✓

Ready:  
 Resume ✓ Cover ✓ Email ✓ LinkedIn draft ✓

[APPLY LINK] [EMAIL RECRUITER]  
 [REFERRAL DRAFT] [OPEN LINKEDIN]  
 [APPROVE] [SKIP]

Your current Telegram service already has approval/skip and external links. Extend it rather than starting a new frontend.

## 16. Cross-link connections with jobs

Accepted LinkedIn connection  
 |  
 v  
 Identify company/person type  
 |  
 +--> Query active jobs for that company  
 |  
 | +--> P0/P1 match -> referral/hiring conversation draft  
 | ^--> no match -> networking-only draft  
 |  
 ^--> store connection as future human-path signal

New P0 job  
 |  
 +--> query saved connections for same company  
 ^--> raise human-path score when a warm route exists

## 17. Gmail reply watcher closes the loop

| Inbound class       | Automatic next step |
|---------------------|---------------------|
| Application receipt | Mark acknowledged   |



| Inbound class      | Automatic next step                                                                 |
|--------------------|-------------------------------------------------------------------------------------|
| Recruiter reply    | Pause follow-ups; Telegram high-priority alert                                      |
| Assessment         | Extract deadline/link; create preparation task                                      |
| Interview invite   | Move stage; prepare interview pack; optionally create calendar event after approval |
| Rejection          | Record stage/reason; stop follow-ups; update analytics                              |
| Offer/compensation | Escalate; never auto-reply                                                          |
| Unknown            | Manual review queue                                                                 |

```
{
  "class": "interview_invite",
  "company": "Acme",
  "role": "MLOps Engineer",
  "confidence": 0.97,
  "recommended_action": "prepare_interview",
  "needs_human": true
}
```

## 18. Follow-ups: automate timing, not spam

Recruiter  
 Day 0: initial outreach  
 Day 4: one follow-up if no reply and role is still live  
 Day 9: final lightweight follow-up if still appropriate  
 Then stop

Employee referral  
 Day 0: referral/advice request  
 Day 5: one polite follow-up  
 Then stop

Generate the follow-up at send time from current state. Never queue a fixed message that might fire after the person already replied.

## 19. Interview mode is automatic

1. Load the exact JD and exact resume version submitted.
2. Research company/product/recent moves using approved sources/connectors.
3. Map likely questions to evidence-bank IDs.
4. Generate a short technical revision plan and company-specific questions.
5. Build/reuse STAR stories from real evidence only.
6. After the interview, capture questions/outcome and update weak areas.

Interview pack

- 1-page company brief
- JD -> evidence matrix
- top technical questions
- behavioral questions + evidence-backed STAR stories
- tell-me-about-yourself answer
- questions to ask interviewer
- 30-minute revision plan
- post-interview thank-you draft

## 20. The database remains the source of truth

jobs(id, company, title, location, url, posted\_at, discovered\_at, fit\_score, human\_path\_score, priority, status, source, jd\_hash)

applications(id, job\_id, applied\_at, resume\_path, cover\_path, status, next\_followup\_at)

contacts(id, company, name, role\_type, title, public\_profile\_url, email, email\_confidence, source)

outreach(id, contact\_id, job\_id, channel, draft, status, sent\_at, message\_id, followup\_due\_at)

connections(id, name, company, title, linkedin\_url, accepted\_at, ...)

events(id, job\_id, source, event\_type, payload\_json, created\_at)

### Why not Notion/Sheets first?

They can be convenient views, but SQLite is better as the transactional source of truth for dedupe, idempotency and background workers. Sync a reporting view to Notion/Sheets later if you want.

## 21. Scheduling: move acquisition to the service that owns it

| Task                      | Best owner                                  | Cadence                            |
|---------------------------|---------------------------------------------|------------------------------------|
| ATS scan                  | Career Ops + local scheduler                | Every 2-4h                         |
| Apify Actor feed          | Apify schedule or Career Ops scheduled scan | Every 2-4h based on cost/freshness |
| Gmail inbox watcher       | Existing Python worker / Gmail API          | Every 5-15m                        |
| Daily Telegram digest     | Python orchestrator                         | 08:00                              |
| Follow-up validity check  | Python orchestrator                         | Daily                              |
| Weekly catch-up/analytics | Python + Career Ops                         | Weekly                             |

If Apify is already running an Actor on a schedule, do not also wake Claude just to launch the same crawl. Let data collection happen where it is cheapest and most reliable; invoke Claude only when there is new data worth reasoning about.

## 22. Permissions and cost guardrails

| Action                                  | Policy                                                      |
|-----------------------------------------|-------------------------------------------------------------|
| Read job/profile files                  | Allow                                                       |
| Write artifacts/logs inside project     | Allow                                                       |
| Run pinned Career Ops commands          | Allow narrow command patterns                               |
| Run Apify Actor                         | Pinned Actor + input limits + maxItems/cost budget          |
| Discover new Actor                      | Human review before production pinning                      |
| Send recruiter email                    | Approval by default; optional strict auto-send policy later |
| Employee/founder outreach               | Always approval                                             |
| LinkedIn send/scrape/browser automation | Deny                                                        |
| Secrets/.env modification               | Deny                                                        |

**Do not use allow-everything**

A connector-first system touches external services. Narrow permissions, daily email caps, Actor input caps, dedupe and audit logs matter more as you automate more.

## 23. Build in four shorter releases

| Release                         | Build                                                                                                  | Done when                                                      |
|---------------------------------|--------------------------------------------------------------------------------------------------------|----------------------------------------------------------------|
| R1 - Connector-fed intelligence | Career Ops native scan + one reviewed Apify provider + score + SQLite + Telegram                       | Fresh ranked jobs arrive automatically without custom scraping |
| R2 - Application factory        | Evidence bank + Career Ops CV/cover + artifact storage                                                 | Every P0 role gets a truthful tailored package                 |
| R3 - Human path + email         | Career Ops contacto + recruiter/employee drafts + Gmail connector/API sending + connection cross-links | Strong jobs include a credible route to a person               |
| R4 - Closed loop                | Gmail reply classifier + follow-ups + Calendar + interview prep + analytics                            | Replies automatically change the pipeline and next actions     |

**What V2 postpones**

Custom dashboards, browser-based auto-apply, custom job scrapers, LinkedIn automation and complicated microservices. They add risk/maintenance before you prove interview conversion.

## 24. Concrete first build sequence

STEP 1 Run Career Ops native scan and store normalized jobs  
 STEP 2 Add one Apify provider entry and verify mapped output  
 STEP 3 Build evidence bank + preferences  
 STEP 4 Add score gate and Telegram digest  
 STEP 5 Generate Career Ops resume/cover only for P0 jobs  
 STEP 6 Add contacto human-path research  
 STEP 7 Upgrade Telegram buttons  
 STEP 8 Connect approved recruiter-email sending  
 STEP 9 Classify Gmail replies and update stages  
 STEP 10 Add interview prep + conversion analytics

**Best MVP demo**

Run one scheduled scan, discover a real fresh Pune/Remote role, produce a score, tailored package, recruiter/contact route and Telegram card. Then simulate/handle a recruiter reply and show the pipeline switching to interview mode.

## 25. Prompts to give Claude Code

- 1) Read Career Ops and my linkedinnotification repo. Design R1 only. Reuse existing plugins/connectors before proposing new integration code.
- 2) Inspect Career Ops plugins/apify. Add one test portals.yml Apify provider for a job source. Do not invent an Actor ID; help me select and review one first.
- 3) Normalize Career Ops scan results into SQLite with URL/JD dedupe and freshness fields. Add tests.
- 4) Build profile/evidence\_bank.yaml from my resume. Mark unsupported or inferred facts NEEDS\_CONFIRMATION.
- 5) Add a structured job scorer using the evidence bank and preferences. Archive jobs below threshold before generating artifacts.

- 6) Extend the existing Telegram approval card with score, freshness, apply link, human-path status, artifact status and safe action buttons.
- 7) Add recruiter email drafting. Keep external sending behind approval until the dedupe/cap/idempotency tests pass.
- 8) Add Gmail reply classification and pipeline state updates. Never auto-reply to offers, compensation or ambiguous messages.

## 26. Test checklist

- Run 20 historical JDs and manually audit score ordering.
- Compare every tailored resume claim against evidence\_bank.yaml.
- Test duplicate URLs, reposted jobs and the same contact across multiple jobs.
- Pin and document the exact Apify Actor used in production; test schema changes/failures.
- Set maxItems/cost caps and confirm an Actor cannot unexpectedly run a huge job.
- Test Telegram double-click/idempotency so one approval cannot send twice.
- Test recruiter reply, rejection, assessment, interview and offer fixtures.
- Kill workers mid-run and verify restart does not duplicate sends/actions.
- Keep LinkedIn automation disabled.

## 27. Metrics: prove that automation is helping

| Metric                             | Use                                              |
|------------------------------------|--------------------------------------------------|
| P0 jobs/week                       | Discovery quality                                |
| P0 -> application                  | Whether your queue is actionable                 |
| Application -> reply               | Targeting/resume quality                         |
| Recruiter outreach -> reply        | Human-path/outreach quality                      |
| Referral ask -> referral           | Warm outreach effectiveness                      |
| Application -> interview           | Primary optimization metric                      |
| Interview -> next round            | Role fit + prep quality                          |
| Source conversion                  | Which ATS/job source is worth crawling           |
| Apify source cost per P0/interview | Whether paid acquisition is worth it             |
| Freshness conversion               | Whether <24h applications outperform older roles |

## 28. Current references

These product capabilities can change. Re-check the current docs before enabling unattended external actions or pinning paid Actors.

- Career Ops repository: <https://github.com/career-ops-hq/career-ops>
- Career Ops Apify plugin skill: <https://github.com/career-ops-hq/career-ops/blob/main/plugins/apify/skill.md>
- Your LinkedIn Notification repository: <https://github.com/srinivassivakumar/linkedinnotification>
- Apify + Claude integration: <https://docs.apify.com/integrations/claude>
- Apify Claude Code CLI integration: <https://docs.apify.com/integrations/claude-code-cli>
- Apify MCP server: <https://docs.apify.com/integrations/mcp>
- Apify Actors: <https://docs.apify.com/actors>
- Claude Google Workspace connectors: <https://support.claude.com/en/articles/10166901-use-google-workspace-connectors>
- Claude Code MCP: <https://code.claude.com/docs/en/mcp>
- Claude Code permissions: <https://code.claude.com/docs/en/permissions>
- Claude Code hooks: <https://code.claude.com/docs/en/hooks-guide>
- LinkedIn automated activity policy: <https://www.linkedin.com/help/linkedin/answer/a1340567>

### Final V2 principle

Use connectors to shorten the plumbing, but keep the career strategy in your own system: verified evidence, ranking, dedupe, contact policy, follow-up rules and conversion analytics. The system should get smaller as it gets more capable.

## Appendix A - Connector decision checklist

| Question                                       | If yes                     | If no                                    |
|------------------------------------------------|----------------------------|------------------------------------------|
| Is there a maintained native connector/plugin? | Use/test it first          | Check official API                       |
| Can it run in the execution mode you need?     | Keep it                    | Use API/thin adapter for background work |
| Does it return structured output?              | Map directly into pipeline | Add normalization step                   |
| Can you cap permissions/cost?                  | Approve for production     | Do not run unattended                    |
| Can you identify data provenance?              | Use for decisions          | Treat as low-confidence/manual review    |
| Does platform ToS allow the automation?        | Proceed within limits      | Keep action manual                       |

## Appendix B - Example daily flow with minimal intervention

02:00 Apify/ATS feed finds new roles  
02:03 Pipeline dedupes + verifies freshness  
02:04 Claude scores only unseen jobs  
02:07 Career Ops builds package for P0 jobs  
02:10 Contacto/public research finds human path  
02:12 Telegram queue updated

08:00 You see:  
3 P0 jobs  
2 recruiter emails ready  
1 employee referral draft  
1 recruiter reply

You spend your time deciding/responding/interviewing,  
not searching, scraping, formatting and tracking.